

```

1 https://www.jianshu.com/p/38286d9859b4
2 Interpreter In Hotspot ( Base OpenJDK 8 )
3
4 _feihui_

```

在阅读这篇文章之前，请各位看官可以先思考一个问题，字节码在 Hotspot 中是编译执行还是解释执行？又或者是？

上面这个问题的答案可以通过在命令行执行 `java -version` 命令得到，得到如下输出：

```

10 openjdk version "1.8.0_222"
11 OpenJDK Runtime Environment (AdoptOpenJDK)(build 1.8.0_222-b10)
12 OpenJDK 64-Bit Server VM (AdoptOpenJDK)(build 25.222-b10, mixed mode)
13 关注最后一行最后的 mixed mode -- 解释执行和编译执行相结合的混合模式 ( 你可以通过关闭 JIT 来达到完全解释执行的效果，你也可以通过开启 AOT 来达到完全编译执行的效果)。

```

今天就先来讲下 Hotspot 与解释执行相关的解释器。

在 OpenJDK 源码(base 8 version 实现)中，可以搜索到两个 Interpreter：

```

19 BytecodeInterpreter.cpp
20 TemplateInterpreter.cpp
21 BytecodeInterpreter: 作为最早的 Interpreter，实现方式非常简单明了，但其执行效率也同样非常"简(bu)单(ren)明(zhi)了(shi)"，以 new 为例看如下代码：

```

```

22
23 switch (opcode) {
24     .....
25     CASE(_new): {
26         u2 index = Bytes::get_Java_u2(pc+1);
27         ConstantPool* constants = istate->method()->constants();
28         if (!constants->tag_at(index).is_unresolved_class()) {
29             // Make sure klass is initialized and doesn't have a finalizer
30             Klass* entry = constants->slot_at(index).get_klass();
31             assert(entry->is_klass(), "Should be resolved klass");
32             Klass* k_entry = (Klass*) entry;
33             assert(k_entry->oop_is_instance(), "Should be InstanceKlass");
34             InstanceKlass* ik = (InstanceKlass*) k_entry;
35             if ( ik->is_initialized() && ik->can_be_fastpath_allocated() ) {
36                 size_t obj_size = ik->size_helper();
37                 oop result = NULL;
38                 // If the TLAB isn't pre-zeroed then we'll have to do it
39                 bool need_zero = !ZeroTLAB;
40                 if (UseTLAB) {
41                     result = (oop) THREAD->tlab().allocate(obj_size);
42                 }
43                 if (result == NULL) {
44                     need_zero = true;
45                     // Try allocate in shared eden
46                 }
47                 retry:
48                 HeapWord* compare_to = *Universe::heap()->top_addr();
49                 HeapWord* new_top = compare_to + obj_size;
50                 if (new_top <= *Universe::heap()->end_addr()) {
51                     if (Atomic::cmpxchg_ptr(new_top, Universe::heap()->top_addr(), compare_to) != compare_to) {
52                         goto retry;
53                     }
54                     result = (oop) compare_to;
55                 }
56                 if (result != NULL) {
57                     // Initialize object (if nonzero size and need) and then the header
58                     if (need_zero) {
59                         HeapWord* to_zero = (HeapWord*) result + sizeof(oopDesc) / oopSize;
60                         obj_size -= sizeof(oopDesc) / oopSize;
61                         if (obj_size > 0) {
62                             memset(to_zero, 0, obj_size * HeapWordSize);
63                         }
64                     }
65                     if (UseBiasedLocking) {
66                         result->set_mark(ik->prototype_header());
67                     } else {
68                         result->set_mark(markOopDesc::prototype());
69                     }
70                     result->set_klass_gap(0);
71                     result->set_klass(k_entry);
72                     SET_STACK_OBJECT(result, 0);
73                     UPDATE_PC_AND_TOS_AND_CONTINUE(3, 1);
74                 }
75             }
76         }
77         // Slow case allocation
78         CALL_VM(InterpreterRuntime::_new(THREAD, METHOD->constants(), index),
79             handle_exception);
80         SET_STACK_OBJECT(THREAD->vm_result(), 0);
81         THREAD->set_vm_result(NULL);
82         UPDATE_PC_AND_TOS_AND_CONTINUE(3, 1);
83     }
84     .....

```

```

85     }
86 可以看到创建一个对象对应如此多的代码( 还不包括其中的宏展开以及方法调用 ), 这效率也能想象得到, 之所以一直保留在源码中, 主要用于方便理解每个字节码的逻辑。
87
88  TemplateInterpreter: 当前默认的 Interpreter, 每一个字节码会有一段对应的精简汇编代码, 同样以 new 为例看如下代码:
89
90  void TemplateTable::initialize() {
91      .....
92      // Java spec bytecodes          ubcp|disp|clvm|iswd  in   out   generator          argument
93      def(Bytecodes::_new              , ubcp|____|clvm|____, vtos, atos, _new              , _              );
94      .....
95  }
96
97  void TemplateTable::_new() {
98      transition(vtos, atos);
99      __ get_unsigned_2_byte_index_at_bcp(rdx, 1);
100     Label slow_case;
101     Label done;
102     Label initialize_header;
103     Label initialize_object; // including clearing the fields
104     Label allocate_shared;
105
106     __ get_cpool_and_tags(rsi, rax);
107     // Make sure the class we're about to instantiate has been resolved.
108     // This is done before loading InstanceKlass to be consistent with the order
109     // how Constant Pool is updated (see ConstantPool::klass_at_put)
110     const int tags_offset = Array<u1>::base_offset_in_bytes();
111     __ cmpb(Address(rax, rdx, Address::times_1, tags_offset),
112            JVM_CONSTANT_Class);
113     __ jcc(Assembler::notEqual, slow_case);
114
115     // get InstanceKlass
116     __ movptr(rsi, Address(rsi, rdx,
117        Address::times_8, sizeof(ConstantPool)));
118
119     // make sure klass is initialized & doesn't have finalizer
120     // make sure klass is fully initialized
121     __ cmpb(Address(rsi,
122        InstanceKlass::init_state_offset()),
123        InstanceKlass::fully_initialized);
124     __ jcc(Assembler::notEqual, slow_case);
125
126     // get instance_size in InstanceKlass (scaled to a count of bytes)
127     __ movl(rdx,
128        Address(rsi,
129        Klass::layout_helper_offset()));
130     // test to see if it has a finalizer or is malformed in some way
131     __ testl(rdx, Klass::_lh_instance_slow_path_bit);
132     __ jcc(Assembler::notZero, slow_case);
133
134     // Allocate the instance
135     // 1) Try to allocate in the TLAB
136     // 2) if fail and the object is large allocate in the shared Eden
137     // 3) if the above fails (or is not applicable), go to a slow case
138     // (creates a new TLAB, etc.)
139
140     const bool allow_shared_alloc =
141         Universe::heap()->supports_inline_contig_alloc() && !CMSIncrementalMode;
142
143     if (UseTLAB) {
144         __ movptr(rax, Address(r15_thread, in_bytes(JavaThread::tlab_top_offset())));
145         __ lea(rbx, Address(rax, rdx, Address::times_1));
146         __ cmpptr(rbx, Address(r15_thread, in_bytes(JavaThread::tlab_end_offset())));
147         __ jcc(Assembler::above, allow_shared_alloc ? allocate_shared : slow_case);
148         __ movptr(Address(r15_thread, in_bytes(JavaThread::tlab_top_offset())), rbx);
149         if (ZeroTLAB) {
150             // the fields have been already cleared
151             __ jmp(initialize_header);
152         } else {
153             // initialize both the header and fields
154             __ jmp(initialize_object);
155         }
156     }
157
158     // Allocation in the shared Eden, if allowed.
159     //
160     // rdx: instance size in bytes
161     if (allow_shared_alloc) {
162         __ bind(allocate_shared);
163
164         ExternalAddress top((address)Universe::heap()->top_addr());
165         ExternalAddress end((address)Universe::heap()->end_addr());
166
167         const Register RtopAddr = rscratch1;
168         const Register RendAddr = rscratch2;
169

```

```

170     __ lea(RtopAddr, top);
171     __ lea(RendAddr, end);
172     __ movptr(rax, Address(RtopAddr, 0));
173
174     // For retries rax gets set by cmpxchgq
175     Label retry;
176     __ bind(retry);
177     __ lea(rbx, Address(rax, rdx, Address::times_1));
178     __ cmpptr(rbx, Address(RendAddr, 0));
179     __ jcc(Assembler::above, slow_case);
180
181     // Compare rax with the top addr, and if still equal, store the new
182     // top addr in rbx at the address of the top addr pointer. Sets ZF if was
183     // equal, and clears it otherwise. Use lock prefix for atomicity on MPs.
184     //
185     // rax: object begin
186     // rbx: object end
187     // rdx: instance size in bytes
188     if (os::is_MP()) {
189         __ lock();
190     }
191     __ cmpxchgptr(rbx, Address(RtopAddr, 0));
192
193     // if someone beat us on the allocation, try again, otherwise continue
194     __ jcc(Assembler::notEqual, retry);
195
196     __ incr_allocated_bytes(r15_thread, rdx, 0);
197 }
198
199 if (UseTLAB || Universe::heap()->supports_inline_contig_alloc()) {
200     // The object is initialized before the header. If the object size is
201     // zero, go directly to the header initialization.
202     __ bind(initialize_object);
203     __ decrementl(rdx, sizeof(oopDesc));
204     __ jcc(Assembler::zero, initialize_header);
205
206     // Initialize object fields
207     __ xorl(rcx, rcx); // use zero reg to clear memory (shorter code)
208     __ shr1(rdx, LogBytesPerLong); // divide by oopSize to simplify the loop
209     {
210         Label loop;
211         __ bind(loop);
212         __ movq(Address(rax, rdx, Address::times_8,
213             sizeof(oopDesc) - oopSize),
214             rcx);
215         __ decrementl(rdx);
216         __ jcc(Assembler::notZero, loop);
217     }
218
219     // initialize object header only.
220     __ bind(initialize_header);
221     if (UseBiasedLocking) {
222         __ movptr(rscratch1, Address(rsi, Klass::prototype_header_offset()));
223         __ movptr(Address(rax, oopDesc::mark_offset_in_bytes()), rscratch1);
224     } else {
225         __ movptr(Address(rax, oopDesc::mark_offset_in_bytes()),
226             (intptr_t) markOopDesc::prototype()); // header (address 0x1)
227     }
228     __ xorl(rcx, rcx); // use zero reg to clear memory (shorter code)
229     __ store_klass_gap(rax, rcx); // zero klass gap for compressed oops
230     __ store_klass(rax, rsi); // store klass last
231
232     {
233         SkipIfEqual skip(_masm, &DTraceAllocProbes, false);
234         // Trigger dtrace event for fastpath
235         __ push(atos); // save the return value
236         __ call_VM_leaf(
237             CAST_FROM_FN_PTR(address, SharedRuntime::dtrace_object_alloc), rax);
238         __ pop(atos); // restore the return value
239     }
240
241     __ jmp(done);
242 }
243
244 // slow case
245 __ bind(slow_case);
246 __ get_constant_pool(c_rarg1);
247 __ get_unsigned_2_byte_index_at_bcp(c_rarg2, 1);
248 call_VM(rax, CAST_FROM_FN_PTR(address, InterpreterRuntime::_new), c_rarg1, c_rarg2);
249 __ verify_oop(rax);
250
251 // continue
252 __ bind(done);
253 }
254
255 上面展示的是生成汇编代码的代码，下面则展示的 Hotspot 初始化后生成的汇编代码：

```

```
256
257 new 187 new [0x000000011828f8c0, 0x000000011828fac0] 512 bytes
258
259 0x000000011828f8c0: push %rax
260 0x000000011828f8c1: jmpq 0x000000011828f8f0
261 0x000000011828f8c6: sub $0x8,%rsp
262 0x000000011828f8ca: vmovss %xmm0, (%rsp)
263 0x000000011828f8cf: jmpq 0x000000011828f8f0
264 0x000000011828f8d4: sub $0x10,%rsp
265 0x000000011828f8d8: vmovsd %xmm0, (%rsp)
266 0x000000011828f8dd: jmpq 0x000000011828f8f0
267 0x000000011828f8e2: sub $0x10,%rsp
268 0x000000011828f8e6: mov %rax, (%rsp)
269 0x000000011828f8ea: jmpq 0x000000011828f8f0
270 0x000000011828f8ef: push %rax
271 0x000000011828f8f0: movzwl 0x1(%r13), %edx
272 0x000000011828f8f5: bswap %edx
273 0x000000011828f8f7: shr $0x10, %edx
274 0x000000011828f8fa: mov -0x18(%rbp), %rsi
275 0x000000011828f8fe: mov 0x8(%rsi), %rsi
276 0x000000011828f902: mov 0x8(%rsi), %rsi
277 0x000000011828f906: mov 0x8(%rsi), %rax
278 0x000000011828f90a: cmpb $0x7, 0x4(%rax, %rdx, 1)
279 0x000000011828f90f: jne 0x000000011828f9e4
280 0x000000011828f915: mov 0x50(%rsi, %rdx, 8), %rsi
281 0x000000011828f91a: cmpb $0x4, 0x162(%rsi)
282 0x000000011828f921: jne 0x000000011828f9e4
283 0x000000011828f927: mov 0x8(%rsi), %edx
284 0x000000011828f92a: test $0x1, %edx
285 0x000000011828f930: jne 0x000000011828f9e4
286 0x000000011828f936: mov 0x60(%r15), %rax
287 0x000000011828f93a: lea (%rax, %rdx, 1), %rbx
288 0x000000011828f93e: cmp 0x70(%r15), %rbx
289 0x000000011828f942: ja 0x000000011828f951
290 0x000000011828f948: mov %rbx, 0x60(%r15)
291 0x000000011828f94c: jmpq 0x000000011828f983
292 0x000000011828f951: movabs $0x7fa4fd6009c8, %r10
293 0x000000011828f95b: movabs $0x7fa4fd6009a0, %r11
294 0x000000011828f965: mov (%r10), %rax
295 0x000000011828f968: lea (%rax, %rdx, 1), %rbx
296 0x000000011828f96c: cmp (%r11), %rbx
297 0x000000011828f96f: ja 0x000000011828f9e4
298 0x000000011828f975: lock cmpxchg %rbx, (%r10)
299 0x000000011828f97a: jne 0x000000011828f968
300 0x000000011828f97c: add %rdx, 0xb8(%r15)
301 0x000000011828f983: sub $0x10, %edx
302 0x000000011828f986: je 0x000000011828f99a
303 0x000000011828f98c: xor %ecx, %ecx
304 0x000000011828f98e: shr $0x3, %edx
305 0x000000011828f991: mov %rcx, 0x8(%rax, %rdx, 8)
306 0x000000011828f996: dec %edx
307 0x000000011828f998: jne 0x000000011828f991
308 0x000000011828f99a: mov 0xa8(%rsi), %r10
309 0x000000011828f9a1: mov %r10, (%rax)
310 0x000000011828f9a4: xor %ecx, %ecx
311 0x000000011828f9a6: mov %rsi, 0x8(%rax)
312 0x000000011828f9aa: cmpb $0x0, -0x87f88f6(%rip) # 0x000000010fa970bb
313 0x000000011828f9b1: je 0x000000011828f9df
314 0x000000011828f9b7: push %rax
315 0x000000011828f9b8: mov %rax, %rdi
316 0x000000011828f9bb: test $0xf, %esp
317 0x000000011828f9c1: je 0x000000011828f9d9
318 0x000000011828f9c7: sub $0x8, %rsp
319 0x000000011828f9cb: callq 0x000000010f6c06c4
320 0x000000011828f9d0: add $0x8, %rsp
321 0x000000011828f9d4: jmpq 0x000000011828f9de
322 0x000000011828f9d9: callq 0x000000010f6c06c4
323 0x000000011828f9de: pop %rax
324 0x000000011828f9df: jmpq 0x000000011828faa8
325 0x000000011828f9e4: mov -0x18(%rbp), %rsi
326 0x000000011828f9e8: mov 0x8(%rsi), %rsi
327 0x000000011828f9ec: mov 0x8(%rsi), %rsi
328 0x000000011828f9f0: movzwl 0x1(%r13), %edx
329 0x000000011828f9f5: bswap %edx
330 0x000000011828f9f7: shr $0x10, %edx
331 0x000000011828f9fa: callq 0x000000011828fa04
332 0x000000011828f9ff: jmpq 0x000000011828faa8
333 0x000000011828fa04: lea 0x8(%rsp), %rax
334 0x000000011828fa09: mov %r13, -0x38(%rbp)
335 0x000000011828fa0d: mov %r15, %rdi
336 0x000000011828fa10: mov %rbp, 0x1d0(%r15)
337 0x000000011828fa17: mov %rax, 0x1c0(%r15)
338 0x000000011828fa1e: test $0xf, %esp
339 0x000000011828fa24: je 0x000000011828fa3c
340 0x000000011828fa2a: sub $0x8, %rsp
341 0x000000011828fa2e: callq 0x000000010f4def58
```

```

342 0x000000011828fa33: add    $0x8,%rsp
343 0x000000011828fa37: jmpq   0x000000011828fa41
344 0x000000011828fa3c: callq  0x000000010f4def58
345 0x000000011828fa41: movabs $0x0,%r10
346 0x000000011828fa4b: mov    %r10,0x1c0(%r15)
347 0x000000011828fa52: movabs $0x0,%r10
348 0x000000011828fa5c: mov    %r10,0x1d0(%r15)
349 0x000000011828fa63: movabs $0x0,%r10
350 0x000000011828fa6d: mov    %r10,0x1c8(%r15)
351 0x000000011828fa74: cmpq   $0x0,0x8(%r15)
352 0x000000011828fa7c: je     0x000000011828fa87
353 0x000000011828fa82: jmpq   0x000000011826a420
354 0x000000011828fa87: mov    0x220(%r15),%rax
355 0x000000011828fa8e: movabs $0x0,%r10
356 0x000000011828fa98: mov    %r10,0x220(%r15)
357 0x000000011828fa9f: mov    -0x38(%rbp),%r13
358 0x000000011828faa3: mov    -0x30(%rbp),%r14
359 0x000000011828faa7: retq
360 0x000000011828faa8: movzbl 0x3(%r13),%ebx
361 0x000000011828faad: add    $0x3,%r13
362 0x000000011828fab1: movabs $0x10faab9e0,%r10
363 0x000000011828fabb: jmpq   *(%r10,%rbx,8)
364 0x000000011828fabf: nop

```

从上面的代码我们可以看出，相对于 BytecodeInterpreter 的实现，TemplateInterpreter 的实现精简了很多。

如何打印字节码对应的汇编代码？示例如下：

```

369 javac -g Solution.java
370 java -XX:-UseCompressedOops -XX:+UnlockDiagnosticVMOptions -XX:+PrintStubCode -XX:+PrintInterpreter -XX:+PrintAssembly Solution
371 // Could not load hsdis-amd64.dylib; library not loadable; PrintAssembly is disabled
372 // mv hsdis-amd64.dylib > /Library/Java/JavaVirtualMachines/jdk1.8.0_221.jdk/Contents/Home/jre/lib
373

```

从 BytecodeInterpreter 到 TemplateInterpreter，性能有了很大的提升（当然提升的方式不仅于此），可能你会觉得手工编写生成汇编的代码维护成本过高（实际上就相当于用汇编编写逻辑，想想都觉得阔怕，曾经也尝试过自动生成，但效果不尽人意），但从这方面去想 -- 本身每个字节码的逻辑基本上是固定不变的，那么维护成本看起来也并不是不可取，况且伊始更关注的执行效率。

Hotspot 为了提升执行效率，除了 Interpreter 的演进，另一个重要的技术则是 JIT -- 通过 profiling 实现运行时优化，这种方法带来的优化使得 JVM 语言在一些场合下可以优于编译执行语言。除了上述两种主要优化，还有诸如 常量替换/循环展开/同步消除/栈上分配/方法内联 等一些列策略性优化。（或许你会问为什么没有提到垃圾收集，就个人而言，我觉得垃圾收集的演进更多的匹配业务场景，例如在后台业务中 Parallel GC 会有更好的吞吐，而在交互业务中 CMS 会有更好的停顿，并不存在一种 GC 适用于所有业务场景）

我们再来简单聊聊 栈顶缓存 -- 将当前操作数据优先存放在栈顶缓存上（寄存器）而非实际栈顶上（内存），我们先通过一个简单例子来了解下。

```

379 int m = a + b
380 // 未使用栈顶缓存，需要七次数据移动
381 a local areas(memory) -> register
382 a register -> stack top(memory)
383 b local areas(memory) -> register
384 b register -> stack top(memory)
385 b stack top(memory) -> register
386 a stack top(memory) -> register
387 m register -> stack top(memory)
388
389 // 使用栈顶缓存，只需要四次数据移动
390 a local areas(memory) -> register
391 a register -> stack top(memory)
392 b local areas(memory) -> register
393 a stack top(memory) -> register

```

从上面的例子我们可以看到，使用栈顶缓存能明显减少一些数据移动。我们以 iconst_0 为例，我们先看来下其模板定义

```

395 def(Bytecodes::iconst_0, __, __, __, __, vtos, itos, iconst, __, 0);
396

```

上面是 iconst_0 的模板定义，vtos 表示这个字节码执行之前栈顶缓存需为空，itos 表示这个字节码执行之后栈顶缓存为int，我们再来看下生成汇编代码部分的代码

```

398
399 void TemplateInterpreterGenerator::set_short_entry_points(Template* t, address& bep, address& cep, address& sep, address& aep, address& iep,
address& lep, address& fep, address& dep, address& vep) {
400     assert(t->is_valid(), "template must exist");
401     switch (t->tos_in()) {
402         case btos:
403         case ctos:
404         case stos:
405             ShouldNotReachHere(); // btos/ctos/stos should use itos.
406             break;
407         case atos: vep = __ pc(); __ pop(atos); aep = __ pc(); generate_and_dispatch(t); break;
408         case itos: vep = __ pc(); __ pop(itos); iep = __ pc(); generate_and_dispatch(t); break;
409         case ltos: vep = __ pc(); __ pop(ltos); lep = __ pc(); generate_and_dispatch(t); break;
410         case ftos: vep = __ pc(); __ pop(ftos); fep = __ pc(); generate_and_dispatch(t); break;
411         case dtos: vep = __ pc(); __ pop(dtos); dep = __ pc(); generate_and_dispatch(t); break;
412         case vtos: set_vtos_entry_points(t, bep, cep, sep, aep, iep, lep, fep, dep, vep); break;
413         default : ShouldNotReachHere(); break;
414     }
415 }
416
417 void TemplateInterpreterGenerator::set_vtos_entry_points(Template* t,
418                                                         address& bep,
419                                                         address& cep,
420                                                         address& sep,

```

```

421                                     address& aep,
422                                     address& iep,
423                                     address& lep,
424                                     address& fep,
425                                     address& dep,
426                                     address& vep) {
427     assert(t->is_valid() && t->tos_in() == vtos, "illegal template");
428     Label L;
429     aep = __ pc(); __ push_ptr(); __ jmp(L);
430     fep = __ pc(); __ push_f(); __ jmp(L);
431     dep = __ pc(); __ push_d(); __ jmp(L);
432     lep = __ pc(); __ push_l(); __ jmp(L);
433     bep = cep = sep =
434     iep = __ pc(); __ push_i();
435     vep = __ pc();
436     __ bind(L);
437     generate_and_dispatch(t);
438 }

```

上面可以看到，当需要栈顶缓存为空时，会首先将缓存数据放入实际栈顶，然后再执行字节码逻辑，如下图汇编代码所示：

```

441 0x000000001184968c0: push    %rax
442 0x000000001184968c1: jmpq    0x000000001184968f0
443 0x000000001184968c6: sub     $0x8,%rsp
444 0x000000001184968ca: vmovss  %xmm0, (%rsp)
445 0x000000001184968cf: jmpq    0x000000001184968f0
446 0x000000001184968d4: sub     $0x10,%rsp
447 0x000000001184968d8: vmovsd  %xmm0, (%rsp)
448 0x000000001184968dd: jmpq    0x000000001184968f0
449 0x000000001184968e2: sub     $0x10,%rsp
450 0x000000001184968e6: mov     %rax, (%rsp)
451 0x000000001184968ea: jmpq    0x000000001184968f0
452 0x000000001184968ef: push    %rax
453 0x000000001184968f0: xor     %eax,%eax
454 0x000000001184968f2: movzbl  0x1(%r13),%ebx
455 0x000000001184968f7: inc     %r13
456 0x000000001184968fa: movabs  $0x10c309160,%r10
457 0x00000000118496904: jmpq    *(%r10,%rbx,8)

```

可以看到 0x000000001184968f0 为 iconst_0 汇编码起始地址，以上的为处理栈顶缓存和跳转逻辑。

最后再来简单聊下 Profiling, Hotspot JIT-compiler 之所以强大的一点，不仅仅因为它能将字节码编辑成机器码（单纯地将字节码编译成机器码并不见得比编译执行语言高效），更重要的结合运行时 Profiling 数据，能将热点路径以最优性能执行 -- 例如：对于一段比较复杂的处理逻辑，有：输入 A 输出 B，而且大部分情况输入都为 A，那么可以直接改优化为比较输入是否为 A，若为 A 则直接输出 B 而不需要执行复杂逻辑，类似于哈夫曼编码。更多详情可阅读郑雨迪博士的专栏《深入拆解Java虚拟机》，专栏干货满满。