

```

// -----
// Generate a native wrapper for a given method. The method takes arguments
// in the Java compiled code convention, marshals them to the native
// convention (handles oops, etc), transitions to native, makes the call,
// returns to java state (possibly blocking), unhandles any result and
// returns.
//
// Critical native functions are a shorthand for the use of
// GetPrimitiveArrayCritical and disallow the use of any other JNI
// functions. The wrapper is expected to unpack the arguments before
// passing them to the callee and perform checks before and after the
// native call to ensure that they GC_locker
// lock_critical/unlock_critical semantics are followed. Some other
// parts of JNI setup are skipped like the tear down of the JNI handle
// block and the check for pending exceptions it's impossible for them
// to be thrown.
//
// They are roughly structured like this:
//   if (GC_locker::needs_gc())
//       SharedRuntime::block_for_jni_critical();
//   transition to thread_in_native
//   unpack array arguments and call native entry point
//   check for safepoint in progress
//   check if any thread suspend flags are set
//   call into JVM and possibly unlock the JNI critical
//   if a GC was suppressed while in the critical native.
//   transition back to thread_in_Java
//   return to caller
//
nmethod* SharedRuntime::generate_native_wrapper(MacroAssembler* masm,
                                                methodHandle method,
                                                int compile_id,
                                                BasicType* in_sig_bt,
                                                VMRegPair* in_regs,
                                                BasicType ret_type) {
    if (method->is_method_handle_intrinsic()) {
        vmIntrinsics::ID iid = method->intrinsic_id();
        intptr_t start = (intptr_t)__ pc();
        int vep_offset = ((intptr_t)__ pc()) - start;
        gen_special_dispatch(masm,
                            method,
                            in_sig_bt,
                            in_regs);
        int frame_complete = ((intptr_t)__ pc()) - start; // not complete, period
        __ flush();
        int stack_slots = SharedRuntime::out_preserve_stack_slots(); // no out slots at all,
        actually
        return nmethod::new_native_nmethod(method,
                                            compile_id,
                                            masm->code(),
                                            vep_offset,
                                            frame_complete,
                                            stack_slots / VMRegImpl::slots_per_word,
                                            in_ByteSize(-1),
                                            in_ByteSize(-1),
                                            (OopMapSet*)NULL);
    }
    bool is_critical_native = true;
    address native_func = method->critical_native_function();
    if (native_func == NULL) {

```

```

    native_func = method->native_function();
    is_critical_native = false;
}
assert(native_func != NULL, "must have function");

// An OopMap for lock (and class if static)
OopMapSet *oop_maps = new OopMapSet();
intptr_t start = (intptr_t) __ pc();

// We have received a description of where all the java arg are located
// on entry to the wrapper. We need to convert these args to where
// the jni function will expect them. To figure out where they go
// we convert the java signature to a C signature by inserting
// the hidden arguments as arg[0] and possibly arg[1] (static method)

const int total_in_args = method->size_of_parameters();
int total_c_args = total_in_args;
if (!is_critical_native) {
    total_c_args += 1;
    if (method->is_static()) {
        total_c_args++;
    }
} else {
    for (int i = 0; i < total_in_args; i++) {
        if (in_sig_bt[i] == T_ARRAY) {
            total_c_args++;
        }
    }
}

BasicType* out_sig_bt = NEW_RESOURCE_ARRAY(BasicType, total_c_args);
VMRegPair* out_regs = NEW_RESOURCE_ARRAY(VMRegPair, total_c_args);
BasicType* in_elem_bt = NULL;

int argc = 0;
if (!is_critical_native) {
    out_sig_bt[argc++] = T_ADDRESS;
    if (method->is_static()) {
        out_sig_bt[argc++] = T_OBJECT;
    }

    for (int i = 0; i < total_in_args ; i++) {
        out_sig_bt[argc++] = in_sig_bt[i];
    }
} else {
    Thread* THREAD = Thread::current();
    in_elem_bt = NEW_RESOURCE_ARRAY(BasicType, total_in_args);
    SignatureStream ss(method->signature());
    for (int i = 0; i < total_in_args ; i++) {
        if (in_sig_bt[i] == T_ARRAY) {
            // Arrays are passed as int, elem* pair
            out_sig_bt[argc++] = T_INT;
            out_sig_bt[argc++] = T_ADDRESS;
            Symbol* atype = ss.as_symbol(CHECK_NULL);
            const char* at = atype->as_C_string();
            if (strlen(at) == 2) {
                assert(at[0] == '[', "must be");
                switch (at[1]) {
                    case 'B': in_elem_bt[i] = T_BYTE; break;
                    case 'C': in_elem_bt[i] = T_CHAR; break;
                    case 'D': in_elem_bt[i] = T_DOUBLE; break;
                    case 'F': in_elem_bt[i] = T_FLOAT; break;
                }
            }
        }
    }
}

```

```

        case 'I': in_elem_bt[i] = T_INT; break;
        case 'J': in_elem_bt[i] = T_LONG; break;
        case 'S': in_elem_bt[i] = T_SHORT; break;
        case 'Z': in_elem_bt[i] = T_BOOLEAN; break;
        default: ShouldNotReachHere();
    }
}
} else {
    out_sig_bt[argc++] = in_sig_bt[i];
    in_elem_bt[i] = T_VOID;
}
if (in_sig_bt[i] != T_VOID) {
    assert(in_sig_bt[i] == ss.type(), "must match");
    ss.next();
}
}
}

// Now figure out where the args must be stored and how much stack space
// they require.
int out_arg_slots;
out_arg_slots = c_calling_convention(out_sig_bt, out_regs, NULL, total_c_args);

// Compute framesize for the wrapper. We need to handle all oops in
// incoming registers

// Calculate the total number of stack slots we will need.

// First count the abi requirement plus all of the outgoing args
int stack_slots = SharedRuntime::out_preserve_stack_slots() + out_arg_slots;

// Now the space for the inbound oop handle area
int total_save_slots = 6 * VMRegImpl::slots_per_word; // 6 arguments passed in registers
if (is_critical_native) {
    // Critical natives may have to call out so they need a save area
    // for register arguments.
    int double_slots = 0;
    int single_slots = 0;
    for (int i = 0; i < total_in_args; i++) {
        if (in_regs[i].first()->is_Register()) {
            const Register reg = in_regs[i].first()->as_Register();
            switch (in_sig_bt[i]) {
                case T_BOOLEAN:
                case T_BYTE:
                case T_SHORT:
                case T_CHAR:
                case T_INT: single_slots++; break;
                case T_ARRAY: // specific to LP64 (7145024)
                case T_LONG: double_slots++; break;
                default: ShouldNotReachHere();
            }
        } else if (in_regs[i].first()->is_XMMRegister()) {
            switch (in_sig_bt[i]) {
                case T_FLOAT: single_slots++; break;
                case T_DOUBLE: double_slots++; break;
                default: ShouldNotReachHere();
            }
        } else if (in_regs[i].first()->is_FloatRegister()) {
            ShouldNotReachHere();
        }
    }
}
}

```

```

    total_save_slots = double_slots * 2 + single_slots;
    // align the save area
    if (double_slots != 0) {
        stack_slots = round_to(stack_slots, 2);
    }
}

int oop_handle_offset = stack_slots;
stack_slots += total_save_slots;

// Now any space we need for handling a klass if static method

int klass_slot_offset = 0;
int klass_offset = -1;
int lock_slot_offset = 0;
bool is_static = false;

if (method->is_static()) {
    klass_slot_offset = stack_slots;
    stack_slots += VMRegImpl::slots_per_word;
    klass_offset = klass_slot_offset * VMRegImpl::stack_slot_size;
    is_static = true;
}

// Plus a lock if needed

if (method->is_synchronized()) {
    lock_slot_offset = stack_slots;
    stack_slots += VMRegImpl::slots_per_word;
}

// Now a place (+2) to save return values or temp during shuffling
// + 4 for return address (which we own) and saved rbp
stack_slots += 6;

// Ok The space we have allocated will look like:
//
//
// FP-> |
//      |-----|
//      | 2 slots for moves |
//      |-----|
//      | lock box (if sync) |
//      |-----| <- lock_slot_offset
//      | klass (if static) |
//      |-----| <- klass_slot_offset
//      | oopHandle area    |
//      |-----| <- oop_handle_offset (6 java arg registers)
//      | outbound memory   |
//      | based arguments   |
//      |
//      |-----|
//      |
// SP-> | out_preserved_slots |
//
//
//
// Now compute actual number of stack words we need rounding to make
// stack properly aligned.
stack_slots = round_to(stack_slots, StackAlignmentInSlots);

```

```

int stack_size = stack_slots * VMRegImpl::stack_slot_size;

// First thing make an ic check to see if we should even be here

// We are free to use all registers as temps without saving them and
// restoring them except rbp. rbp is the only callee save register
// as far as the interpreter and the compiler(s) are concerned.

const Register ic_reg = rax;
const Register receiver = j_rarg0;

Label hit;
Label exception_pending;

assert_different_registers(ic_reg, receiver, rscratch1);
__ verify_oop(receiver);
__ load_klass(rscratch1, receiver);
__ cmpq(ic_reg, rscratch1);
__ jcc(Assembler::equal, hit);

__ jump(RuntimeAddress(SharedRuntime::get_ic_miss_stub()));

// Verified entry point must be aligned
__ align(8);

__ bind(hit);

int vep_offset = ((intptr_t)__ pc()) - start;

// The instruction at the verified entry point must be 5 bytes or longer
// because it can be patched on the fly by make_non_entrant. The stack bang
// instruction fits that requirement.

// Generate stack overflow check

if (UseStackBanging) {
    __ bang_stack_with_offset(StackShadowPages*os::vm_page_size());
} else {
    // need a 5 byte instruction to allow MT safe patching to non-entrant
    __ fat_nop();
}

// Generate a new frame for the wrapper.
__ enter();
// -2 because return address is already present and so is saved rbp
__ subptr(rsp, stack_size - 2*wordSize);

// Frame is now completed as far as size and linkage.
int frame_complete = ((intptr_t)__ pc()) - start;

if (UseRTMLocking) {
    // Abort RTM transaction before calling JNI
    // because critical section will be large and will be
    // aborted anyway. Also nmethod could be deoptimized.
    __ xabort(0);
}

#ifdef ASSERT
{
    Label L;
    __ mov(rax, rsp);
    __ andptr(rax, -16); // must be 16 byte boundary (see amd64 ABI)

```

```

    __ cmpptr(rax, rsp);
    __ jcc(Assembler::equal, L);
    __ stop("improperly aligned stack");
    __ bind(L);
}
#endif /* ASSERT */

// We use r14 as the oop handle for the receiver/klass
// It is callee save so it survives the call to native

const Register oop_handle_reg = r14;

if (is_critical_native) {
    check_needs_gc_for_critical_native(masm, stack_slots, total_c_args, total_in_args,
                                       oop_handle_offset, oop_maps, in_regs, in_sig_bt);
}

//
// We immediately shuffle the arguments so that any vm call we have to
// make from here on out (sync slow path, jvmti, etc.) we will have
// captured the oops from our caller and have a valid oopMap for
// them.

// -----
// The Grand Shuffle

// The Java calling convention is either equal (linux) or denser (win64) than the
// c calling convention. However the because of the jni_env argument the c calling
// convention always has at least one more (and two for static) arguments than Java.
// Therefore if we move the args from java -> c backwards then we will never have
// a register->register conflict and we don't have to build a dependency graph
// and figure out how to break any cycles.
//

// Record esp-based slot for receiver on stack for non-static methods
int receiver_offset = -1;

// This is a trick. We double the stack slots so we can claim
// the oops in the caller's frame. Since we are sure to have
// more args than the caller doubling is enough to make
// sure we can capture all the incoming oop args from the
// caller.
//
OopMap* map = new OopMap(stack_slots * 2, 0 /* arg_slots*/);

// Mark location of rbp (someday)
// map->set_callee_saved(VMRegImpl::stack2reg( stack_slots - 2), stack_slots * 2, 0,
vmreg(rbp));

// Use eax, ebx as temporaries during any memory-memory moves we have to do
// All inbound args are referenced based on rbp and all outbound args via rsp.

#ifdef ASSERT
    bool reg_destroyed[RegisterImpl::number_of_registers];
    bool freg_destroyed[XMMRegisterImpl::number_of_registers];
    for ( int r = 0 ; r < RegisterImpl::number_of_registers ; r++ ) {
        reg_destroyed[r] = false;
    }
    for ( int f = 0 ; f < XMMRegisterImpl::number_of_registers ; f++ ) {
        freg_destroyed[f] = false;
    }

```

```

#endif /* ASSERT */

// This may iterate in two different directions depending on the
// kind of native it is. The reason is that for regular JNI natives
// the incoming and outgoing registers are offset upwards and for
// critical natives they are offset down.
GrowableArray<int> arg_order(2 * total_in_args);
VMRegPair tmp_vmreg;
tmp_vmreg.set1(rbx->as_VMReg());

if (!is_critical_native) {
    for (int i = total_in_args - 1, c_arg = total_c_args - 1; i >= 0; i--, c_arg--) {
        arg_order.push(i);
        arg_order.push(c_arg);
    }
} else {
    // Compute a valid move order, using tmp_vmreg to break any cycles
    ComputeMoveOrder cmo(total_in_args, in_regs, total_c_args, out_regs, in_sig_bt,
arg_order, tmp_vmreg);
}

int temploc = -1;
for (int ai = 0; ai < arg_order.length(); ai += 2) {
    int i = arg_order.at(ai);
    int c_arg = arg_order.at(ai + 1);
    __ block_comment(err_msg("move %d -> %d", i, c_arg));
    if (c_arg == -1) {
        assert(is_critical_native, "should only be required for critical natives");
        // This arg needs to be moved to a temporary
        __ mov(tmp_vmreg.first()->as_Register(), in_regs[i].first()->as_Register());
        in_regs[i] = tmp_vmreg;
        temploc = i;
        continue;
    } else if (i == -1) {
        assert(is_critical_native, "should only be required for critical natives");
        // Read from the temporary location
        assert(temploc != -1, "must be valid");
        i = temploc;
        temploc = -1;
    }
}
#ifdef ASSERT
    if (in_regs[i].first()->is_Register()) {
        assert(!reg_destroyed[in_regs[i].first()->as_Register()->encoding()], "destroyed
reg!");
    } else if (in_regs[i].first()->is_XMMRegister()) {
        assert(!freg_destroyed[in_regs[i].first()->as_XMMRegister()->encoding()], "destroyed
reg!");
    }
    if (out_regs[c_arg].first()->is_Register()) {
        reg_destroyed[out_regs[c_arg].first()->as_Register()->encoding()] = true;
    } else if (out_regs[c_arg].first()->is_XMMRegister()) {
        freg_destroyed[out_regs[c_arg].first()->as_XMMRegister()->encoding()] = true;
    }
}
#endif /* ASSERT */
switch (in_sig_bt[i]) {
    case T_ARRAY:
        if (is_critical_native) {
            unpack_array_argument(masm, in_regs[i], in_elem_bt[i], out_regs[c_arg + 1],
out_regs[c_arg]);
            c_arg++;
        }
}
#ifdef ASSERT

```

```

        if (out_regs[c_arg].first()->is_Register()) {
            reg_destroyed[out_regs[c_arg].first()->as_Register()->encoding()] = true;
        } else if (out_regs[c_arg].first()->is_XMMRegister()) {
            freg_destroyed[out_regs[c_arg].first()->as_XMMRegister()->encoding()] = true;
        }
    #endif
        break;
    }
    case T_OBJECT:
        assert(!is_critical_native, "no oop arguments");
        object_move(masm, map, oop_handle_offset, stack_slots, in_regs[i], out_regs[c_arg],
                    ((i == 0) && (!is_static)),
                    &receiver_offset);

        break;
    case T_VOID:
        break;

    case T_FLOAT:
        float_move(masm, in_regs[i], out_regs[c_arg]);
        break;

    case T_DOUBLE:
        assert( i + 1 < total_in_args &&
                in_sig_bt[i + 1] == T_VOID &&
                out_sig_bt[c_arg+1] == T_VOID, "bad arg list");
        double_move(masm, in_regs[i], out_regs[c_arg]);
        break;

    case T_LONG :
        long_move(masm, in_regs[i], out_regs[c_arg]);
        break;

    case T_ADDRESS: assert(false, "found T_ADDRESS in java args");

    default:
        move32_64(masm, in_regs[i], out_regs[c_arg]);
    }
}

int c_arg;

// Pre-load a static method's oop into r14. Used both by locking code and
// the normal JNI call code.
if (!is_critical_native) {
    // point c_arg at the first arg that is already loaded in case we
    // need to spill before we call out
    c_arg = total_c_args - total_in_args;

    if (method->is_static()) {

        // load oop into a register
        __ movoop(oop_handle_reg, JNIHandles::make_local(method->method_holder()-
>java_mirror()));

        // Now handle the static class mirror it's known not-null.
        __ movptr(Address(rsp, klass_offset), oop_handle_reg);
        map->set_oop(VMRegImpl::stack2reg(klass_slot_offset));

        // Now get the handle
        __ lea(oop_handle_reg, Address(rsp, klass_offset));
        // store the klass handle as second argument
        __ movptr(c_rarg1, oop_handle_reg);

```



```

    // and protect the arg if we must spill
    c_arg--;
}
} else {
    // For JNI critical methods we need to save all registers in save_args.
    c_arg = 0;
}

// Change state to native (we save the return address in the thread, since it might not
// be pushed on the stack when we do a a stack traversal). It is enough that the pc()
// points into the right code segment. It does not have to be the correct return pc.
// We use the same pc/oopMap repeatedly when we call out

intptr_t the_pc = (intptr_t) __ pc();
oop_maps->add_gc_map(the_pc - start, map);

__ set_last_Java_frame(rsp, noreg, (address)the_pc);

// We have all of the arguments setup at this point. We must not touch any register
// argument registers at this point (what if we save/restore them there are no oop?

{
    SkipIfEqual skip(masm, &DTraceMethodProbes, false);
    // protect the args we've loaded
    save_args(masm, total_c_args, c_arg, out_regs);
    __ mov_metadata(c_rarg1, method());
    __ call_VM_leaf(
        CAST_FROM_FN_PTR(address, SharedRuntime::dtrace_method_entry),
        r15_thread, c_rarg1);
    restore_args(masm, total_c_args, c_arg, out_regs);
}

// RedefineClasses() tracing support for obsolete method entry
if (RC_TRACE_IN_RANGE(0x00001000, 0x00002000)) {
    // protect the args we've loaded
    save_args(masm, total_c_args, c_arg, out_regs);
    __ mov_metadata(c_rarg1, method());
    __ call_VM_leaf(
        CAST_FROM_FN_PTR(address, SharedRuntime::rc_trace_method_entry),
        r15_thread, c_rarg1);
    restore_args(masm, total_c_args, c_arg, out_regs);
}

// Lock a synchronized method

// Register definitions used by locking and unlocking

const Register swap_reg = rax; // Must use rax for cpxchg instruction
const Register obj_reg = rbx; // Will contain the oop
const Register lock_reg = r13; // Address of compiler lock object (BasicLock)
const Register old_hdr = r13; // value of old header at unlock time

Label slow_path_lock;
Label lock_done;

if (method->is_synchronized()) {
    assert(!is_critical_native, "unhandled");

    const int mark_word_offset = BasicLock::displaced_header_offset_in_bytes();

    // Get the handle (the 2nd argument)

```

```

__ mov(oop_handle_reg, c_rarg1);

// Get address of the box

__ lea(lock_reg, Address(rsp, lock_slot_offset * VMRegImpl::stack_slot_size));

// Load the oop from the handle
__ movptr(obj_reg, Address(oop_handle_reg, 0));

if (UseBiasedLocking) {
    __ biased_locking_enter(lock_reg, obj_reg, swap_reg, rscratch1, false, lock_done,
&slow_path_lock);
}

// Load immediate 1 into swap_reg %rax
__ movl(swap_reg, 1);

// Load (object->mark() | 1) into swap_reg %rax
__ orptr(swap_reg, Address(obj_reg, 0));

// Save (object->mark() | 1) into BasicLock's displaced header
__ movptr(Address(lock_reg, mark_word_offset), swap_reg);

if (os::is_MP()) {
    __ lock();
}

// src -> dest iff dest == rax else rax <- dest
__ cmpxchgptr(lock_reg, Address(obj_reg, 0));
__ jcc(Assembler::equal, lock_done);

// Hmm should this move to the slow path code area???

// Test if the oopMark is an obvious stack pointer, i.e.,
// 1) (mark & 3) == 0, and
// 2) rsp <= mark < mark + os::pagesize()
// These 3 tests can be done by evaluating the following
// expression: ((mark - rsp) & (3 - os::vm_page_size())),
// assuming both stack pointer and pagesize have their
// least significant 2 bits clear.
// NOTE: the oopMark is in swap_reg %rax as the result of cmpxchg

__ subptr(swap_reg, rsp);
__ andptr(swap_reg, 3 - os::vm_page_size());

// Save the test result, for recursive case, the result is zero
__ movptr(Address(lock_reg, mark_word_offset), swap_reg);
__ jcc(Assembler::notEqual, slow_path_lock);

// Slow path will re-enter here

__ bind(lock_done);
}

// Finally just about ready to make the JNI call

// get JNIEnv* which is first argument to native
if (!is_critical_native) {
    __ lea(c_rarg0, Address(r15_thread, in_bytes(JavaThread::jni_environment_offset())));
}

```

```

// Now set thread in native
__ movl(Address(r15_thread, JavaThread::thread_state_offset()), _thread_in_native);

__ call(RuntimeAddress(native_func));

// Verify or restore cpu control state after JNI call
__ restore_cpu_control_state_after_jni();

// Unpack native results.
switch (ret_type) {
case T_BOOLEAN: __ c2bool(rax);          break;
case T_CHAR   : __ movzwl(rax, rax);     break;
case T_BYTE   : __ sign_extend_byte(rax); break;
case T_SHORT  : __ sign_extend_short(rax); break;
case T_INT    : /* nothing to do */      break;
case T_DOUBLE :
case T_FLOAT  :
    // Result is in xmm0 we'll save as needed
    break;
case T_ARRAY:      // Really a handle
case T_OBJECT:     // Really a handle
    break; // can't de-handle until after safepoint check
case T_VOID: break;
case T_LONG: break;
default          : ShouldNotReachHere();
}

// Switch thread to "native transition" state before reading the synchronization state.
// This additional state is necessary because reading and testing the synchronization
// state is not atomic w.r.t. GC, as this scenario demonstrates:
//   Java thread A, in _thread_in_native state, loads _not_synchronized and is preempted.
//   VM thread changes sync state to synchronizing and suspends threads for GC.
//   Thread A is resumed to finish this native method, but doesn't block here since it
//   didn't see any synchronization is progress, and escapes.
__ movl(Address(r15_thread, JavaThread::thread_state_offset()), _thread_in_native_trans);

if(os::is_MP()) {
    if (UseMembar) {
        // Force this write out before the read below
        __ membar(Assembler::Membar_mask_bits(
            Assembler::LoadLoad | Assembler::LoadStore |
            Assembler::StoreLoad | Assembler::StoreStore));
    } else {
        // Write serialization page so VM thread can do a pseudo remote membar.
        // We use the current thread pointer to calculate a thread specific
        // offset to write to within the page. This minimizes bus traffic
        // due to cache line collision.
        __ serialize_memory(r15_thread, rcx);
    }
}

Label after_transition;

// check for safepoint operation in progress and/or pending suspend requests
{
    Label Continue;

    __ cmp32(ExternalAddress((address)SafepointSynchronize::address_of_state()),
        SafepointSynchronize::_not_synchronized);

    Label L;
    __ jcc(Assembler::notEqual, L);
}

```

```

__ cmpl(Address(r15_thread, JavaThread::suspend_flags_offset()), 0);
__ jcc(Assembler::equal, Continue);
__ bind(L);

// Don't use call_VM as it will see a possible pending exception and forward it
// and never return here preventing us from clearing _last_native_pc down below.
// Also can't use call_VM_leaf either as it will check to see if rsi & rdi are
// preserved and correspond to the bcp/locals pointers. So we do a runtime call
// by hand.
//
save_native_result(masm, ret_type, stack_slots);
__ mov(c_rarg0, r15_thread);
__ mov(r12, rsp); // remember sp
__ subptr(rsp, frame::arg_reg_save_area_bytes); // windows
__ andptr(rsp, -16); // align stack as required by ABI
if (!is_critical_native) {
__ call(RuntimeAddress(CAST_FROM_FN_PTR(address,
JavaThread::check_special_condition_for_native_trans)));
} else {
__ call(RuntimeAddress(CAST_FROM_FN_PTR(address,
JavaThread::check_special_condition_for_native_trans_and_transition)));
}
__ mov(rsp, r12); // restore sp
__ reinit_heapbase();
// Restore any method result value
restore_native_result(masm, ret_type, stack_slots);

if (is_critical_native) {
// The call above performed the transition to thread_in_Java so
// skip the transition logic below.
__ jmpb(after_transition);
}

__ bind(Continue);
}

// change thread state
__ movl(Address(r15_thread, JavaThread::thread_state_offset()), _thread_in_Java);
__ bind(after_transition);

Label reguard;
Label reguard_done;
__ cmpl(Address(r15_thread, JavaThread::stack_guard_state_offset()),
JavaThread::stack_guard_yellow_disabled);
__ jcc(Assembler::equal, reguard);
__ bind(reguard_done);

// native result if any is live

// Unlock
Label unlock_done;
Label slow_path_unlock;
if (method->is_synchronized()) {

// Get locked oop from the handle we passed to jni
__ movptr(obj_reg, Address(oop_handle_reg, 0));

Label done;

if (UseBiasedLocking) {
__ biased_locking_exit(obj_reg, old_hdr, done);
}
}

```

```

// Simple recursive lock?

__ cmpptr(Address(rsp, lock_slot_offset * VMRegImpl::stack_slot_size),
(int32_t)NULL_WORD);
__ jcc(Assembler::equal, done);

// Must save rax if it is live now because cmpxchg must use it
if (ret_type != T_FLOAT && ret_type != T_DOUBLE && ret_type != T_VOID) {
    save_native_result(masm, ret_type, stack_slots);
}

// get address of the stack lock
__ lea(rax, Address(rsp, lock_slot_offset * VMRegImpl::stack_slot_size));
// get old displaced header
__ movptr(old_hdr, Address(rax, 0));

// Atomic swap old header if oop still contains the stack lock
if (os::is_MP()) {
    __ lock();
}
__ cmpxchgptr(old_hdr, Address(obj_reg, 0));
__ jcc(Assembler::notEqual, slow_path_unlock);

// slow path re-enters here
__ bind(unlock_done);
if (ret_type != T_FLOAT && ret_type != T_DOUBLE && ret_type != T_VOID) {
    restore_native_result(masm, ret_type, stack_slots);
}

__ bind(done);
}
{
    SkipIfEqual skip(masm, &DTraceMethodProbes, false);
    save_native_result(masm, ret_type, stack_slots);
    __ mov_metadata(c_rarg1, method());
    __ call_VM_leaf(
        CAST_FROM_FN_PTR(address, SharedRuntime::dtrace_method_exit),
        r15_thread, c_rarg1);
    restore_native_result(masm, ret_type, stack_slots);
}

__ reset_last_Java_frame(false, true);

// Unpack oop result
if (ret_type == T_OBJECT || ret_type == T_ARRAY) {
    Label L;
    __ testptr(rax, rax);
    __ jcc(Assembler::zero, L);
    __ movptr(rax, Address(rax, 0));
    __ bind(L);
    __ verify_oop(rax);
}

if (!is_critical_native) {
    // reset handle block
    __ movptr(rcx, Address(r15_thread, JavaThread::active_handles_offset()));
    __ movl(Address(rcx, JNIHandleBlock::top_offset_in_bytes()), (int32_t)NULL_WORD);
}

// pop our frame

```

```

__ leave();

if (!is_critical_native) {
    // Any exception pending?
    __ cmpptr(Address(r15_thread, in_bytes(Thread::pending_exception_offset())),
(int32_t)NULL_WORD);
    __ jcc(Assembler::notEqual, exception_pending);
}

// Return

__ ret(0);

// Unexpected paths are out of line and go here

if (!is_critical_native) {
    // forward the exception
    __ bind(exception_pending);

    // and forward the exception
    __ jump(RuntimeAddress(StubRoutines::forward_exception_entry()));
}

// Slow path locking & unlocking
if (method->is_synchronized()) {

    // BEGIN Slow path lock
    __ bind(slow_path_lock);

    // has last_Java_frame setup. No exceptions so do vanilla call not call_VM
    // args are (oop obj, BasicLock* lock, JavaThread* thread)

    // protect the args we've loaded
    save_args(masm, total_c_args, c_arg, out_regs);

    __ mov(c_rarg0, obj_reg);
    __ mov(c_rarg1, lock_reg);
    __ mov(c_rarg2, r15_thread);

    // Not a leaf but we have last_Java_frame setup as we want
    __ call_VM_leaf(CAST_FROM_FN_PTR(address, SharedRuntime::complete_monitor_locking_C), 3);
    restore_args(masm, total_c_args, c_arg, out_regs);

#ifdef ASSERT
    { Label L;
        __ cmpptr(Address(r15_thread, in_bytes(Thread::pending_exception_offset())),
(int32_t)NULL_WORD);
        __ jcc(Assembler::equal, L);
        __ stop("no pending exception allowed on exit from monitorenter");
        __ bind(L);
    }
#endif
    __ jmp(lock_done);

    // END Slow path lock

    // BEGIN Slow path unlock
    __ bind(slow_path_unlock);

    // If we haven't already saved the native result we must save it now as xmm registers
    // are still exposed.

    if (ret_type == T_FLOAT || ret_type == T_DOUBLE ) {

```

```

    save_native_result(masm, ret_type, stack_slots);
}

__ lea(c_rarg1, Address(rsp, lock_slot_offset * VMRegImpl::stack_slot_size));

__ mov(c_rarg0, obj_reg);
__ mov(r12, rsp); // remember sp
__ subptr(rsp, frame::arg_reg_save_area_bytes); // windows
__ andptr(rsp, -16); // align stack as required by ABI

// Save pending exception around call to VM (which contains an EXCEPTION_MARK)
// NOTE that obj_reg == rbx currently
__ movptr(rbx, Address(r15_thread, in_bytes(Thread::pending_exception_offset())));
__ movptr(Address(r15_thread, in_bytes(Thread::pending_exception_offset())),
(int32_t)NULL_WORD);

__ call(RuntimeAddress(CAST_FROM_FN_PTR(address,
SharedRuntime::complete_monitor_unlocking_C)));
__ mov(rsp, r12); // restore sp
__ reinit_heapbase();
#ifdef ASSERT
{
    Label L;
    __ cmpptr(Address(r15_thread, in_bytes(Thread::pending_exception_offset())),
(int)NULL_WORD);
    __ jcc(Assembler::equal, L);
    __ stop("no pending exception allowed on exit complete_monitor_unlocking_C");
    __ bind(L);
}
#endif /* ASSERT */

__ movptr(Address(r15_thread, in_bytes(Thread::pending_exception_offset())), rbx);

if (ret_type == T_FLOAT || ret_type == T_DOUBLE ) {
    restore_native_result(masm, ret_type, stack_slots);
}
__ jmp(unlock_done);

// END Slow path unlock

} // synchronized

// SLOW PATH Reguard the stack if needed

__ bind(reguard);
save_native_result(masm, ret_type, stack_slots);
__ mov(r12, rsp); // remember sp
__ subptr(rsp, frame::arg_reg_save_area_bytes); // windows
__ andptr(rsp, -16); // align stack as required by ABI
__ call(RuntimeAddress(CAST_FROM_FN_PTR(address, SharedRuntime::reguard_yellow_pages)));
__ mov(rsp, r12); // restore sp
__ reinit_heapbase();
restore_native_result(masm, ret_type, stack_slots);
// and continue
__ jmp(reguard_done);

__ flush();

nmethod *nm = nmethod::new_native_nmethod(method,
                                           compile_id,
                                           masm->code(),

```

```
        vep_offset,  
        frame_complete,  
        stack_slots / VMRegImpl::slots_per_word,  
        (is_static ? in_ByteSize(klass_offset) :  
in_ByteSize(receiver_offset)),  
in_ByteSize(lock_slot_offset*VMRegImpl::stack_slot_size),  
        oop_maps);  
  
    if (is_critical_native) {  
        nm->set_lazy_critical_native(true);  
    }  
  
    return nm;  
}
```